Full-Time Job

Principal Engineer, Core ML Infra, Training and Experimentation - Mountain View, San Francisco or Sunnyvale

About this role

Full-Time Job (Level 8)

Location: Mountain View, San Francisco, Sunnyvale

Job family: Software Engineering

Job subfamily: Software Engineering

Product Area: TI Development

Updated Oct 2 | 162 views

Team members

Meenu Gaba ↗
Mountain View
Sr. Director of Engineering, Core ML (hiring manager)

Kevin Hicar ↗
New York
Executive Recruiter, Principal/Distinguished…

Express Interest ⍰

Use the field below to send a message to this role's contact person. Your message is private: only that contact and the staffing team can see it. If you're in Tech, the hiring manager may request your GRAD from People Ops (see go/sharemygrad ↗).

If you have questions about the role, but aren't yet re to express interest, email the hiring manager listed the left. If you're transferring into Google from anoth Bet, check out go/bet-mobility-faq ↗ first.

For more information regarding short term assignments, including eligibility, please visit go/transfers ↗.

Have non-role specific questions about the Leaders Transfer Process at Google? Thinking about a Trans with Level change? Curious about ways to optimize application? Sign up for a confidential, 1:1 mobility session at go/leadershipmobility-oh ↗ or review curated resources at go/leadershipmobility ↗ .

To: **Kevin Hicar** ↗
New York
Executive Recruiter, Principal/Distingui

Transfer message*

EXPRESS INTERE

Description

Core ML is a sub team in the Google AI and Infrastructure org which is responsible for delivering world-class compute, platforms, and tools powering Google's AI-first products, research, and cloud at scale, price, and efficiency. Core ML's mission is to drive ML excellence for Google and the world. The org strives to make it easier to create a well lit and efficient path for AI experimentation, development, and productionization, and align infrastructure initiatives with new areas of research and product innovation. This helps us better meet the challenge of the rapidly evolving hardware and software space around AI. As the Core ML Training & Experimentation Principal Engineer, you will lead design and development of the Core ML training, bulk inference,. and experimentation infrastructure. You will drive the technical strategy and road map across the training ecosystem for Core ML. The Principal Engineer will engage closely with researchers, applied ML leads, and product UTLs to design and develop the world's best ML bulk inference and experimentation ecosystem, extending Google's leadership position in AI. They will also drive ML velocity efforts across Google. This is a highly impactful and strategic position requiring significant technical expertise and influencing skills to innovate and get endorsement from multiple domain experts and senior executives.

Minimum Qualifications:

- Bachelor's degree in Computer Science, or similar technical field of study or equivalent
- 15 years of professional experience as a hands-on software engineer, or 13 years with a Bachelor's degree
- Experience with ML technologies and infrastructure with regard to training, bulk inference, and experimentation, inclusive of long running workflows
- A demonstrated track record of technical innovation with one or more relevant domains (such as AI/ML) in a large, complex, and cross functional enterprise
- Acumen in Machine Learning coupled with exceptionally strong API and systems design skills

Preferred Qualifications:

- An advanced degree in Computer Science, or similar technical field of study or equivalent
- 20 years of professional experience, or 15 years with an advanced degree
- Strong understanding of the needs of advanced ML developers
- Expertise in reasoning and quantifying relative impacts and risks of technical work that could require the involvement of thousands of engineers
- Deep technical expertise in systems and software with the leadership skills required to influence technical leaders across the company
- Ability to quickly ramp up in new subject areas with short notice to guide a few weeks of mission-critical research to applied infrastructure

Responsibilities:

- Provide leadership in designing and implementing user-friendly, high-performance experimentation and bulk inference stack and ecosystem for research and Google products
- Work cross-functionally with Google Research and products to design developer solutions to streamline the research to production path for large models and drive ML velocity to enable quick iteration cycles for research and product teams
- Develop a structure of other technical leads in the area, both by defining technical goals and orienting teams around decisions
- Provide leadership in firefighting situations or situations that require complex and immediate triaging, technical problems, and other emergency work and helping managers of associated teams understand the larger technical issues
- Act as a senior reviewer or go-to person for large scale changes or technical decision-making, participate in calibration and promotion feedback, and develop/drive appropriate decision-making and technical review processes

Please note:

This role is eligible for a transfer with a level change (TLC). To learn about Googler eligibility and process requirements for applying and potentially transferring into this role from one level below, see go/twlc ↗ and general transfer criteria ↗.

- When submitting your application, please email your last two GRAD reviews (PDF) to Kevin Michael Hicar.
- We will review all applications and reach out as there is interest to move forward.
- Please do not reach out directly to the hiring manager for a coffee or info chat, as Google has moved away from pre-interview conversations to help remove bias, one way or another, from the interview and qualification process.

The US base salary range for this full-time position is $294,000-$414,000 + bonus + equity + benefits. Transfer compensation is determined algorithmically and is non-negotiable. Your recruiter will share more about the specific salary for your targeted location during the hiring process. Learn more about how a transfer may affect your compensation package ↗, how location changes affect compensation ↗, and about benefits at Google at go/benefits ↗.

Skills: [Google Infrastructure]   [Inference]   [API Development]   [Data Modeling]   (Machine Learning)

◯ My skills   ◯ Other skills